

AI HYPE, HOPE OR HORROR?

INTRO TO AI SAFETY



IMAGINE...
AI FOR SCAMS



CHENGCHENG TAN

- UCLA Ling & CS
- Stanford MSCS
- FAR.AI Communications
- W2D2 Board



* all views my own, here in **personal** capacity today

INTRO TO SAFETY

- What's the Problem?
- Risks
- Approaches

Resources

WHAT'S THE PROBLEM?

INTRO TO AI SAFETY



**We must take the
risks of AI as seriously
as other major global
challenges.**

Demis Hassabis
Google DeepMind Co-Founder & CEO
Nobel Prize for Protein Folding

It's kind of weird to
think that what you do
might kill everyone,
but still do it.

Sam Altman
OpenAI CEO



A large great white shark swims gracefully through a dense school of smaller fish in the deep blue ocean. The shark's body is a dark grey-blue on top, fading to a lighter shade on its underside. Its mouth is slightly open, showing its rows of sharp teeth. The surrounding water is filled with the silhouettes of many smaller fish, creating a sense of depth and movement.

AI safety threats are
overhyped B.S.

**'Godfather of AI'
leaves Google and
warns of dangers ahead**

Geoffrey Hinton
Nobel Laureate for Neural Networks
Univ of Toronto Professor Emeritus



STATEMENT OF RISK

Center for AI Safety statement signed by

600+

AI experts & public figures

CEOs of OpenAI, DeepMind, Anthropic

STATEMENT OF RISK

**Mitigating the risk of extinction
from AI should be a global priority
alongside other societal-scale
risks such as pandemics and
nuclear war.**

No, ChatGPT can't kill us...
at least not today.



AGI

Artificial
GENERAL
Intelligence



Ultimate goal is
Superintelligence...



Risk Factors: Time Horizon & Speed
Speed of capabilities is improving
very fast.



Can AI scaling continue
through **2030**?

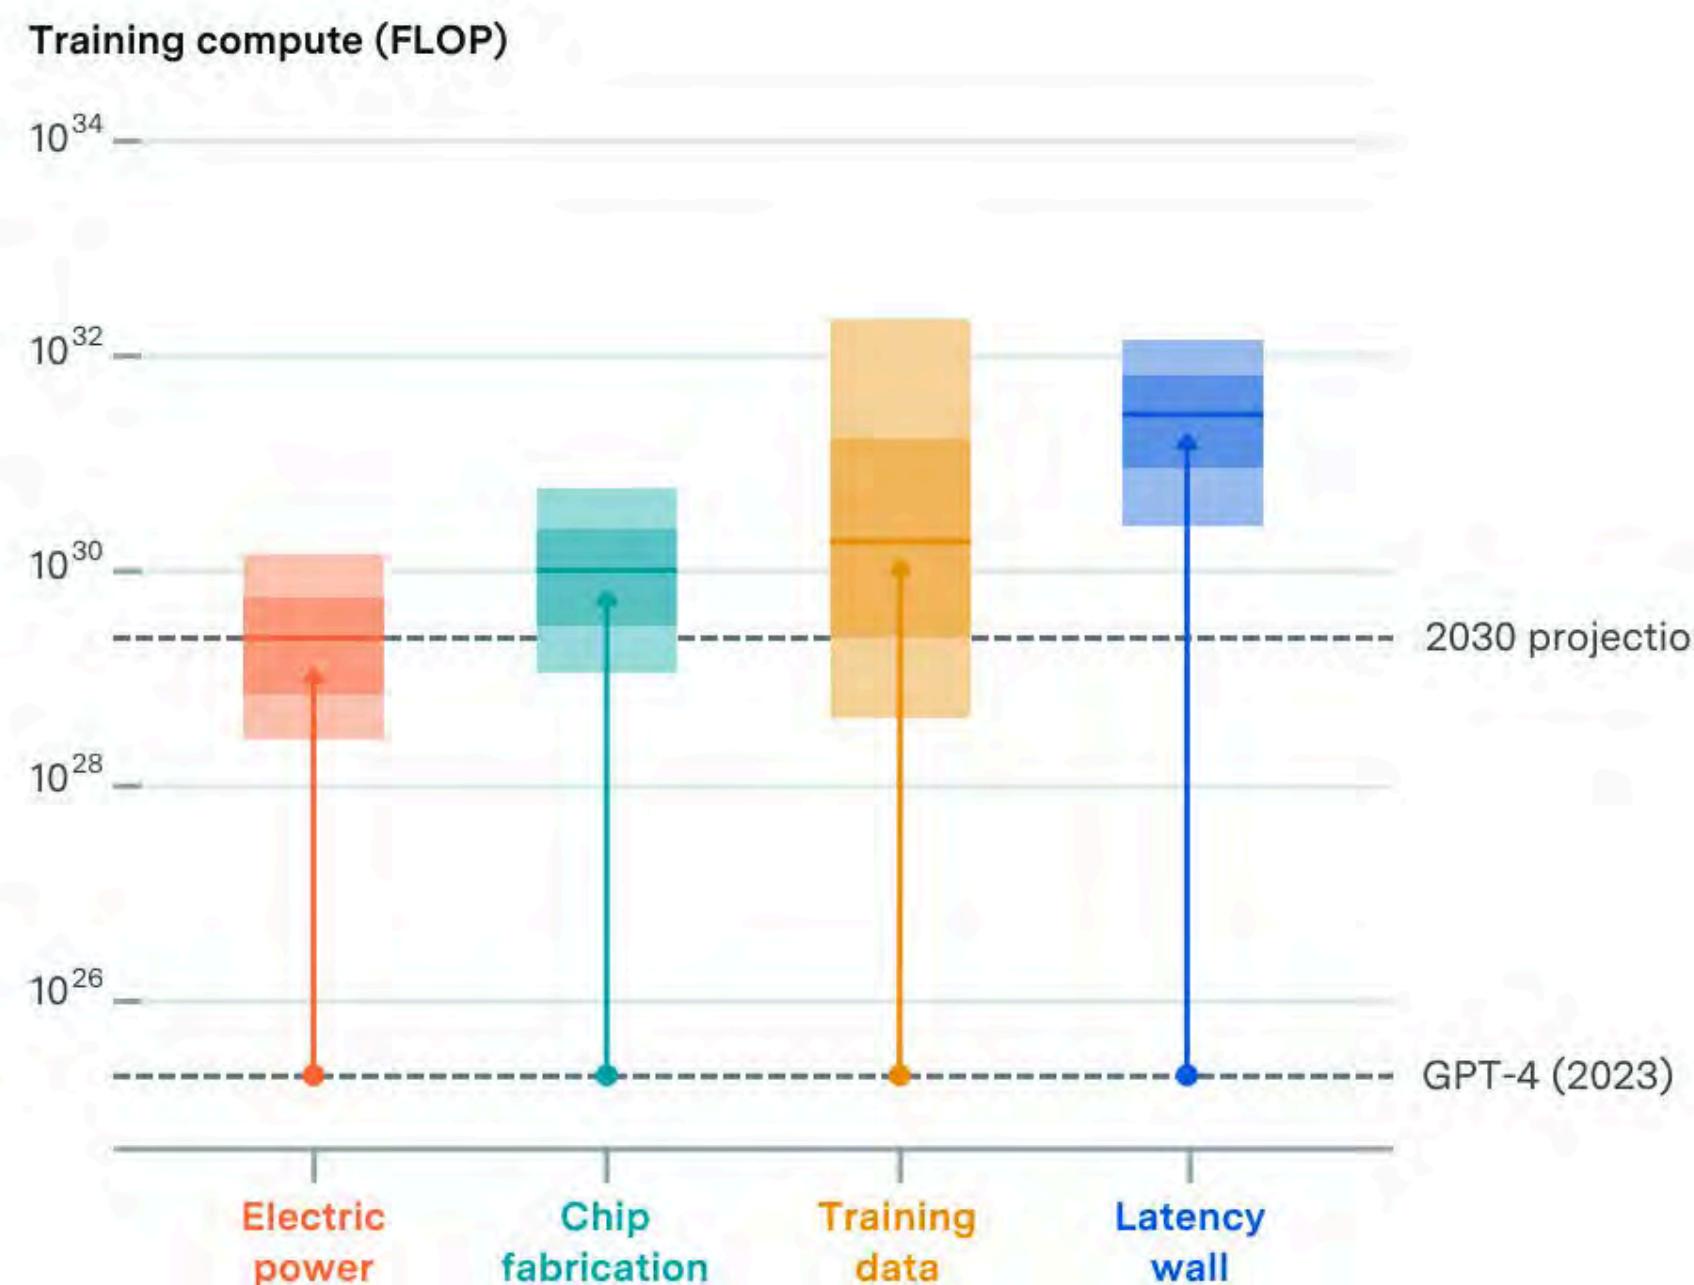
Source: Can AI Scaling Continue Through 2030?

A leap as large as from **GPT-2 to GPT-4**
is on trend by 2030.



Source: Can AI Scaling Continue Through 2030?

Despite challenges, AI growth can likely continue through **2030**.



Source: Can AI Scaling Continue Through 2030?

Possible Risks

1. Misuse by Humans
2. Societal Destabilization
3. Misalignment





MIS-Use by Humans

- Intentional
Malicious Use
- Unintentional
Accidents

Societal Destabilization

- Misinformation & Deep Fakes
- Privacy
- Unemployment
- Superhuman Persuasion

Alignment: AI Goals to Human Values

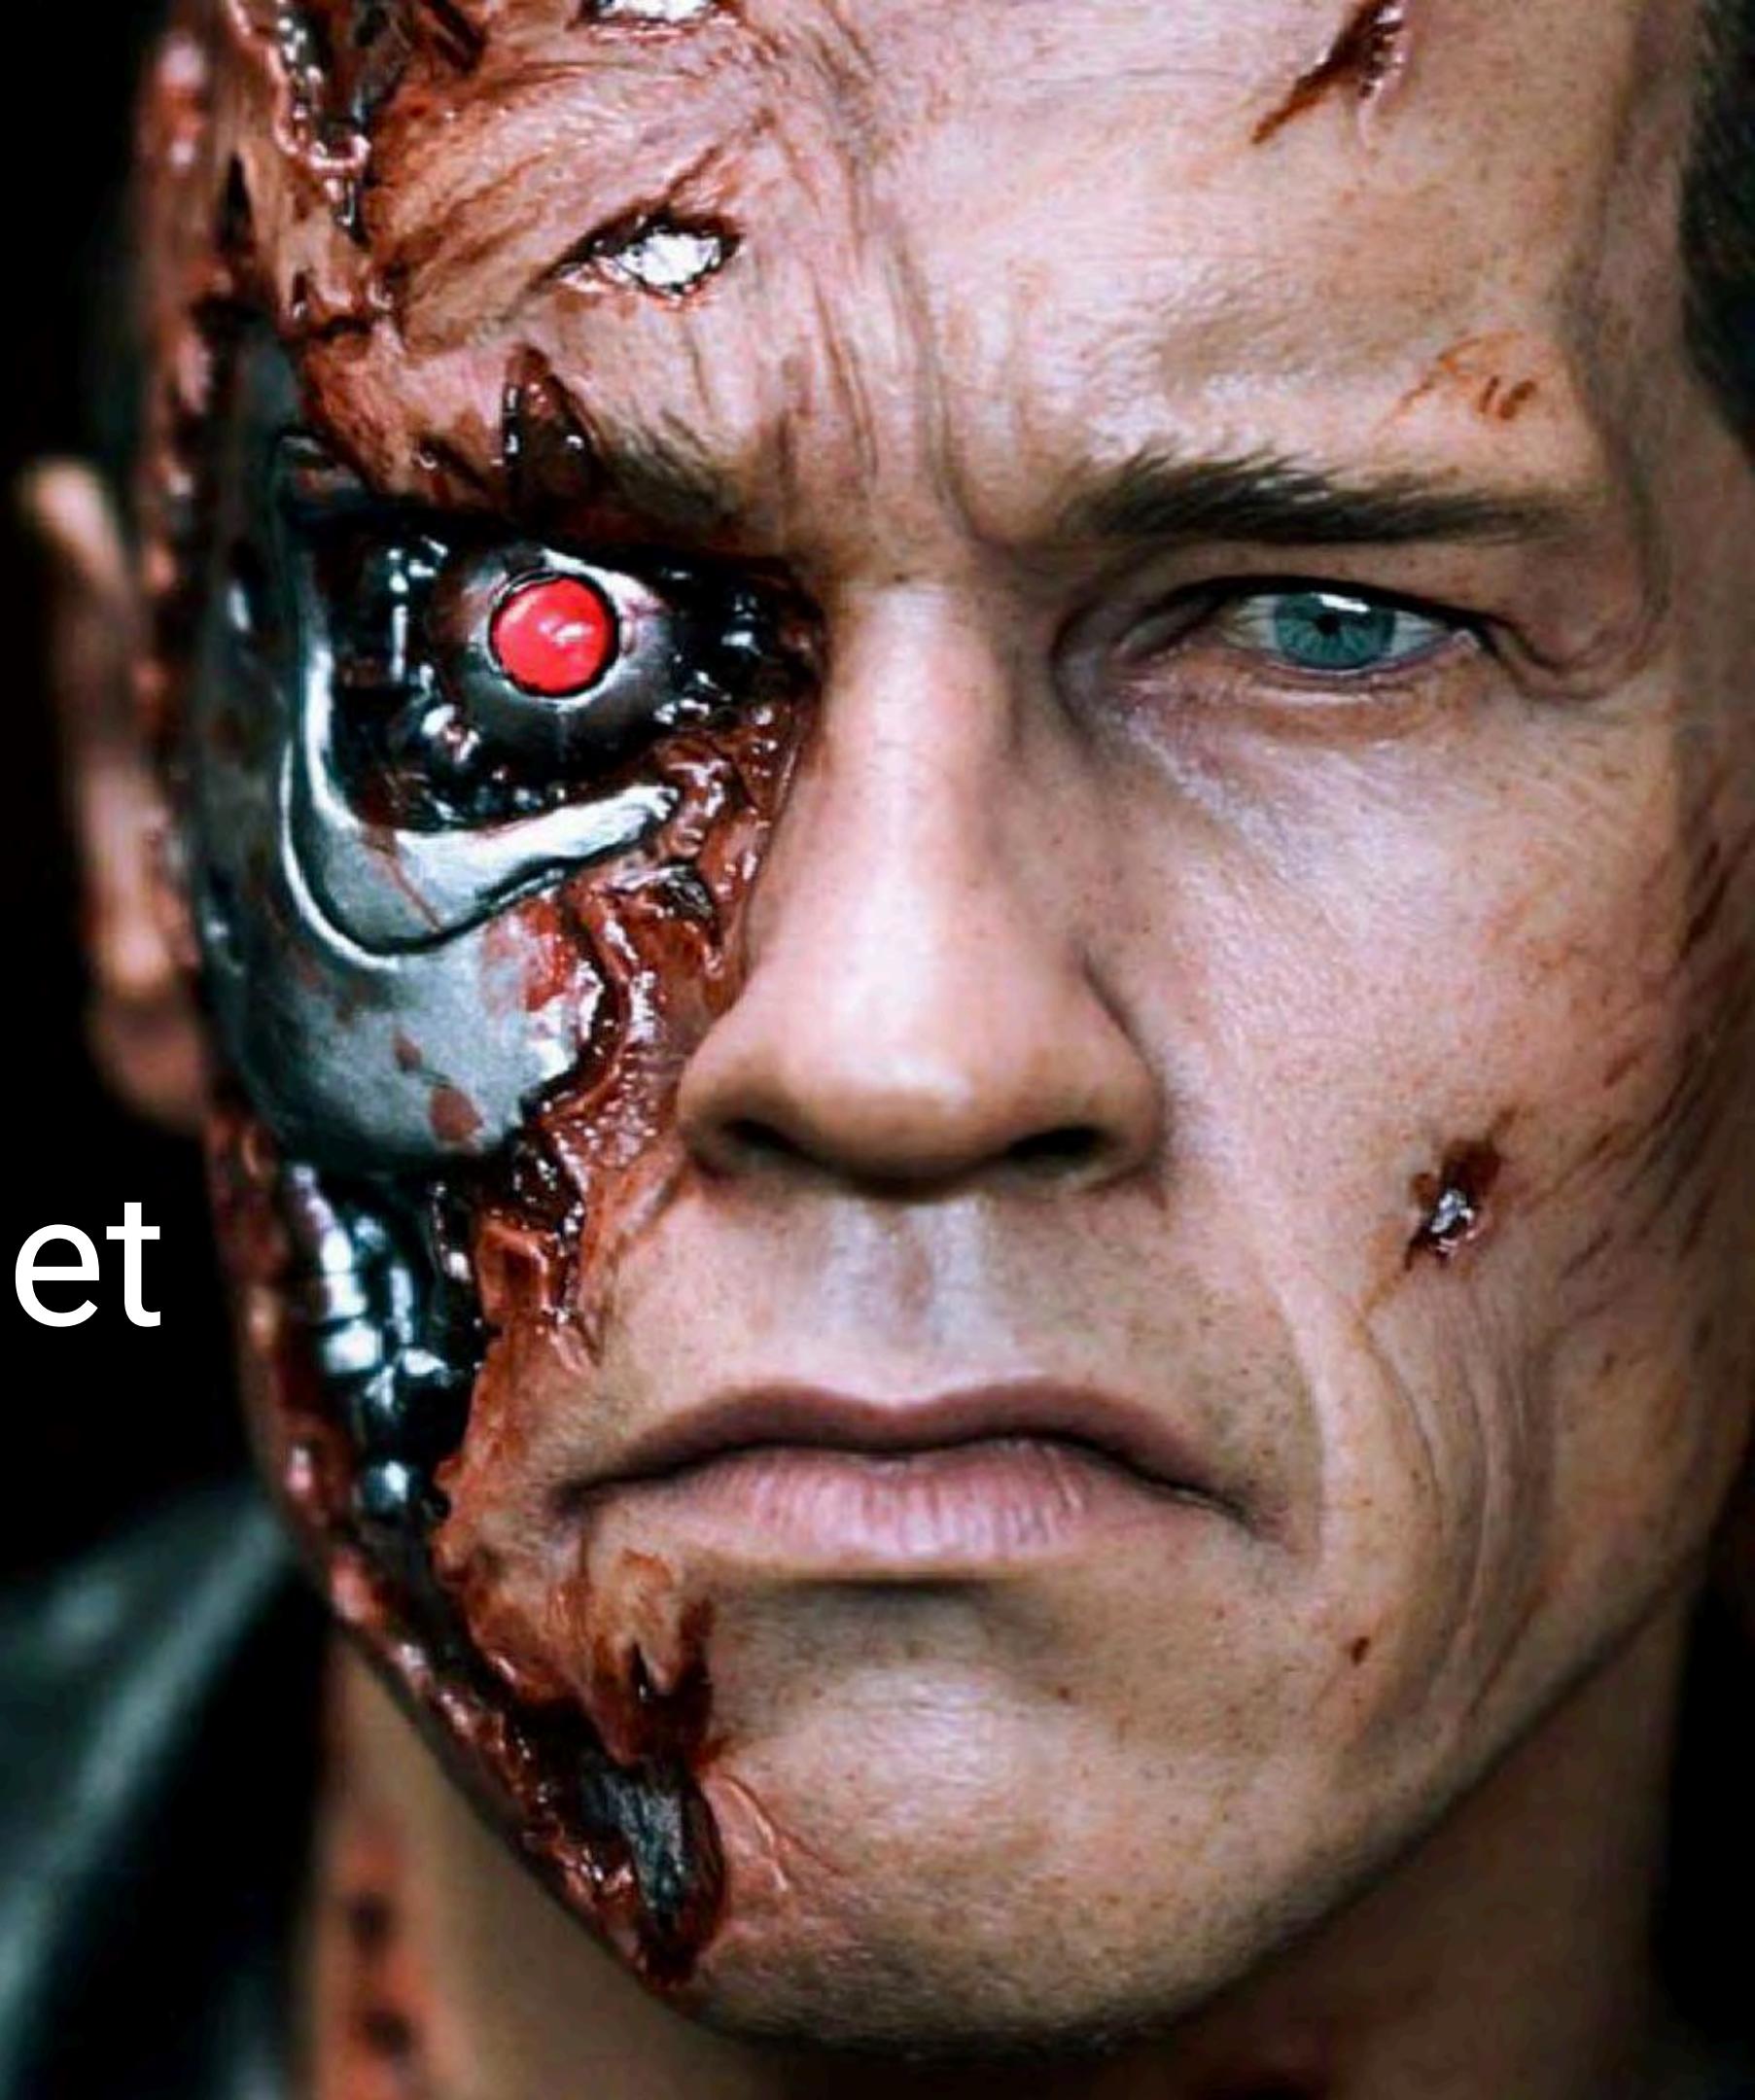


MIS-Alignment

Power Seeking

- Rogue AI, e.g.,

Terminator & Skynet



MIS-Alignment

Power Seeking

- Rogue AI, e.g.,
Terminator & Skynet
- Paperclip Maximizer &

Gorilla Problem





Misalignment today

AI Ethics: Fairness, Bias, Privacy

WHAT ARE SOME APPROACHES?

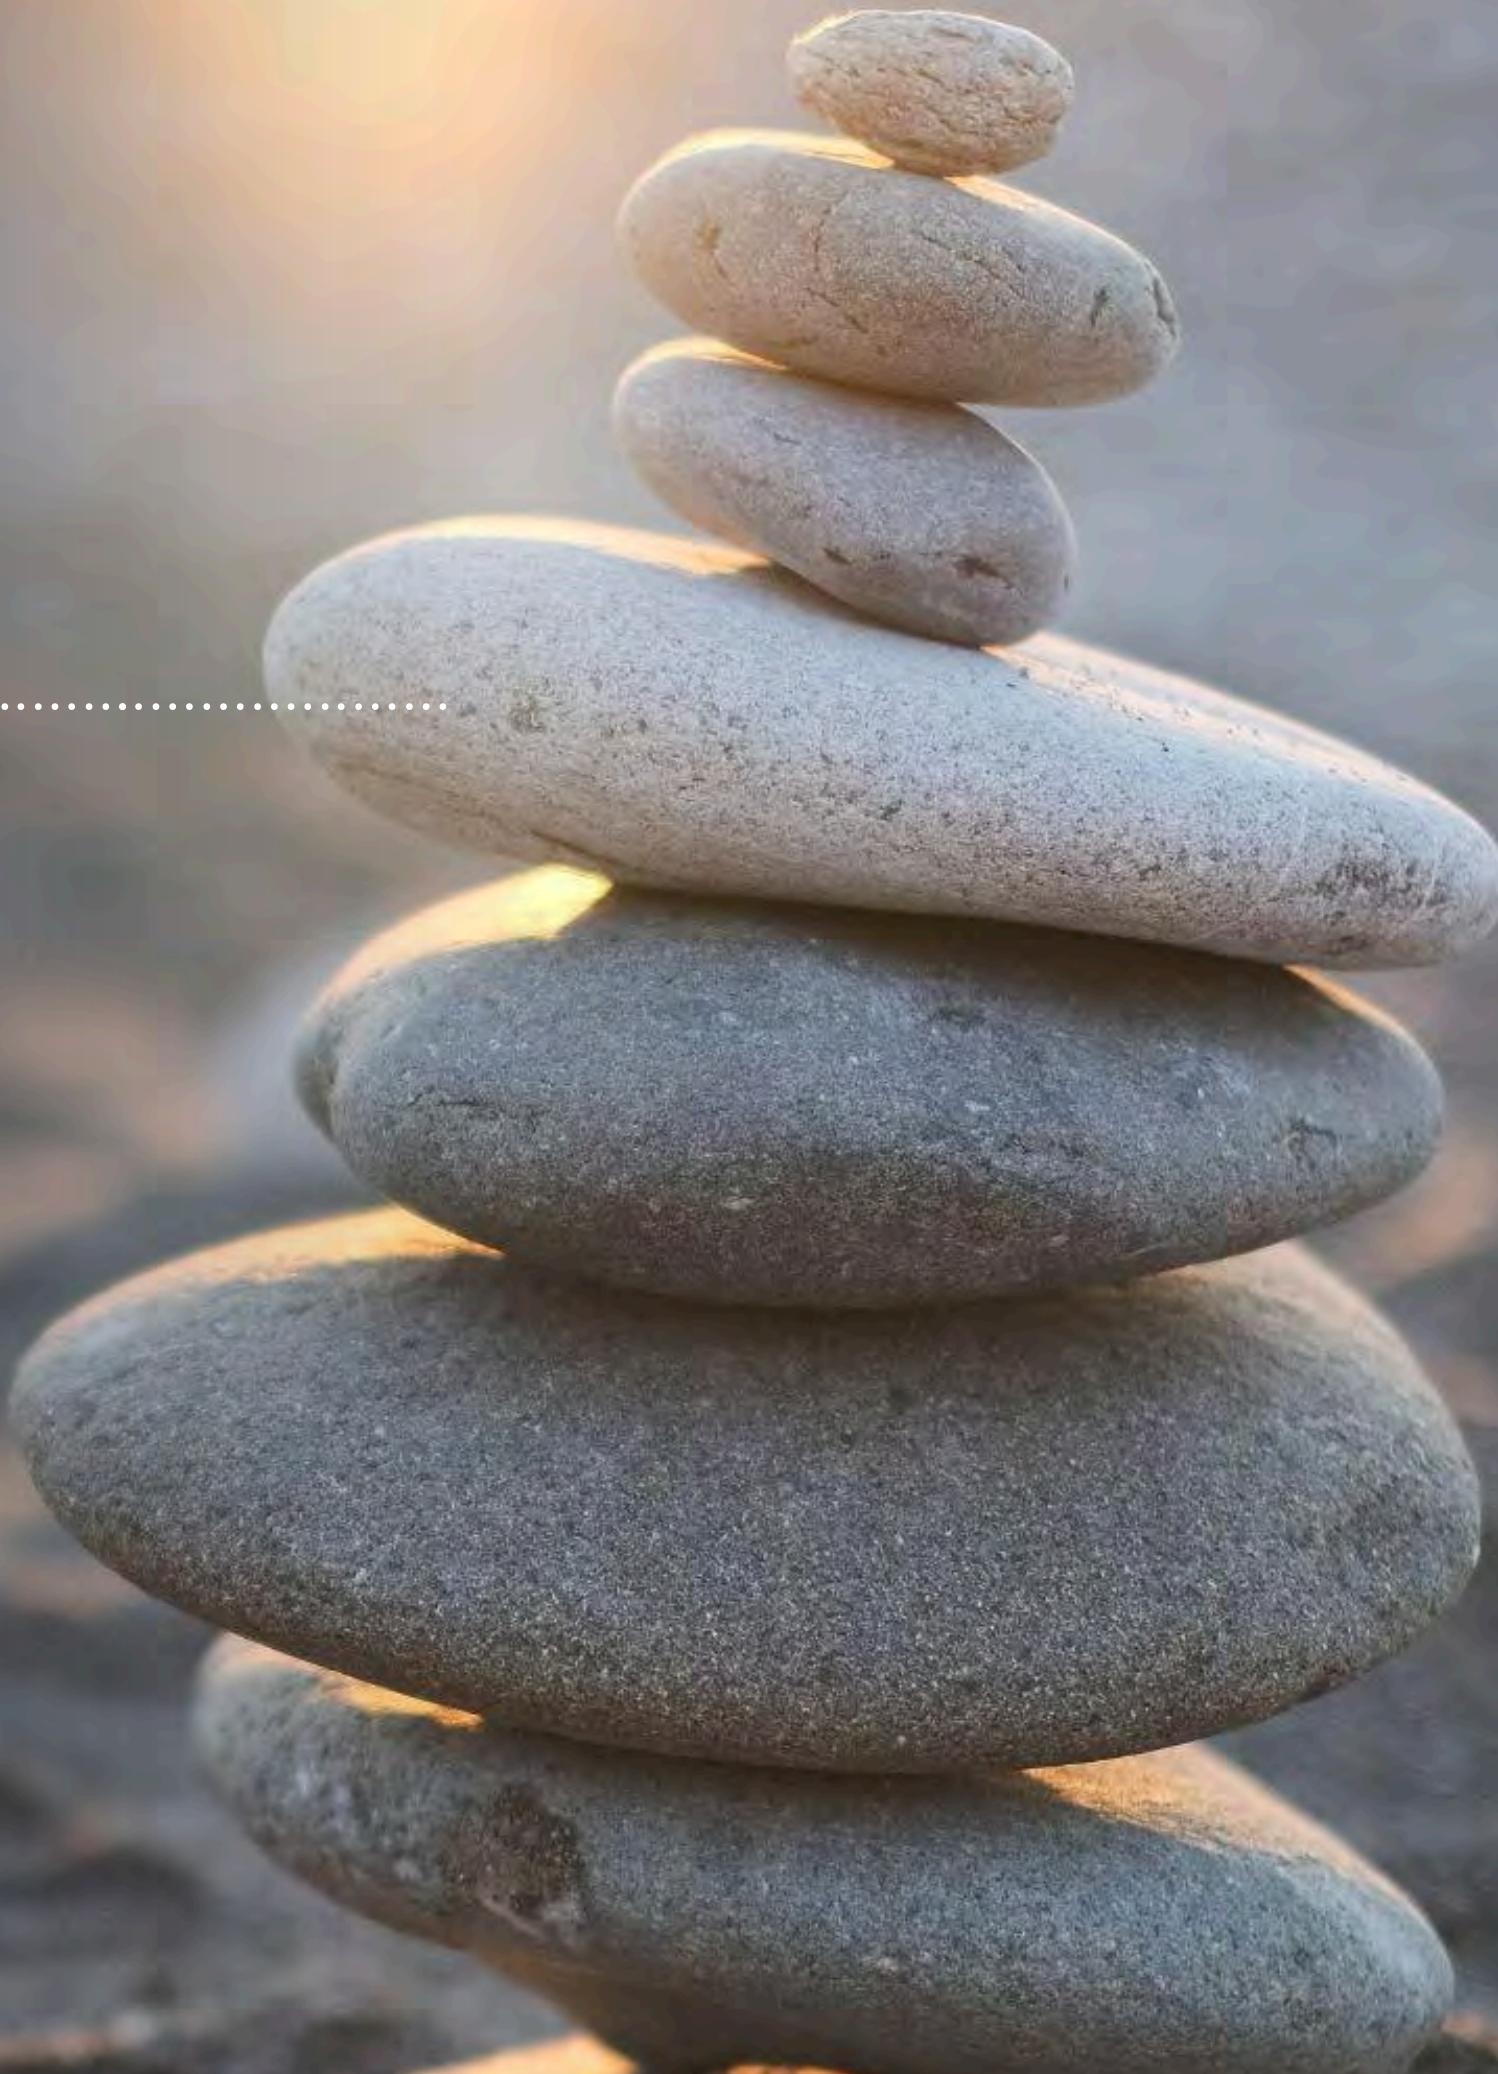
INTRO TO AI SAFETY

CONCEPTS

- Value Alignment
- Evaluations & Robustness
- Scalable Oversight
- Interpretability
- Governance

VALUE ALIGNMENT

Align AI goals with
human values
Ex: GPT-3 Alignment



**GPT-3 is like
Shoggoth**

RLHF:
Conversations +
Instructions



EVALUATIONS & ROBUSTNESS

Test AI for reliability
& resilience
Ex: Jailbreak, Red-teaming

Frontier AIs can be
easily manipulated
into helping with
any harmful request.



Vulnerable models include those from:

Google

OpenAI

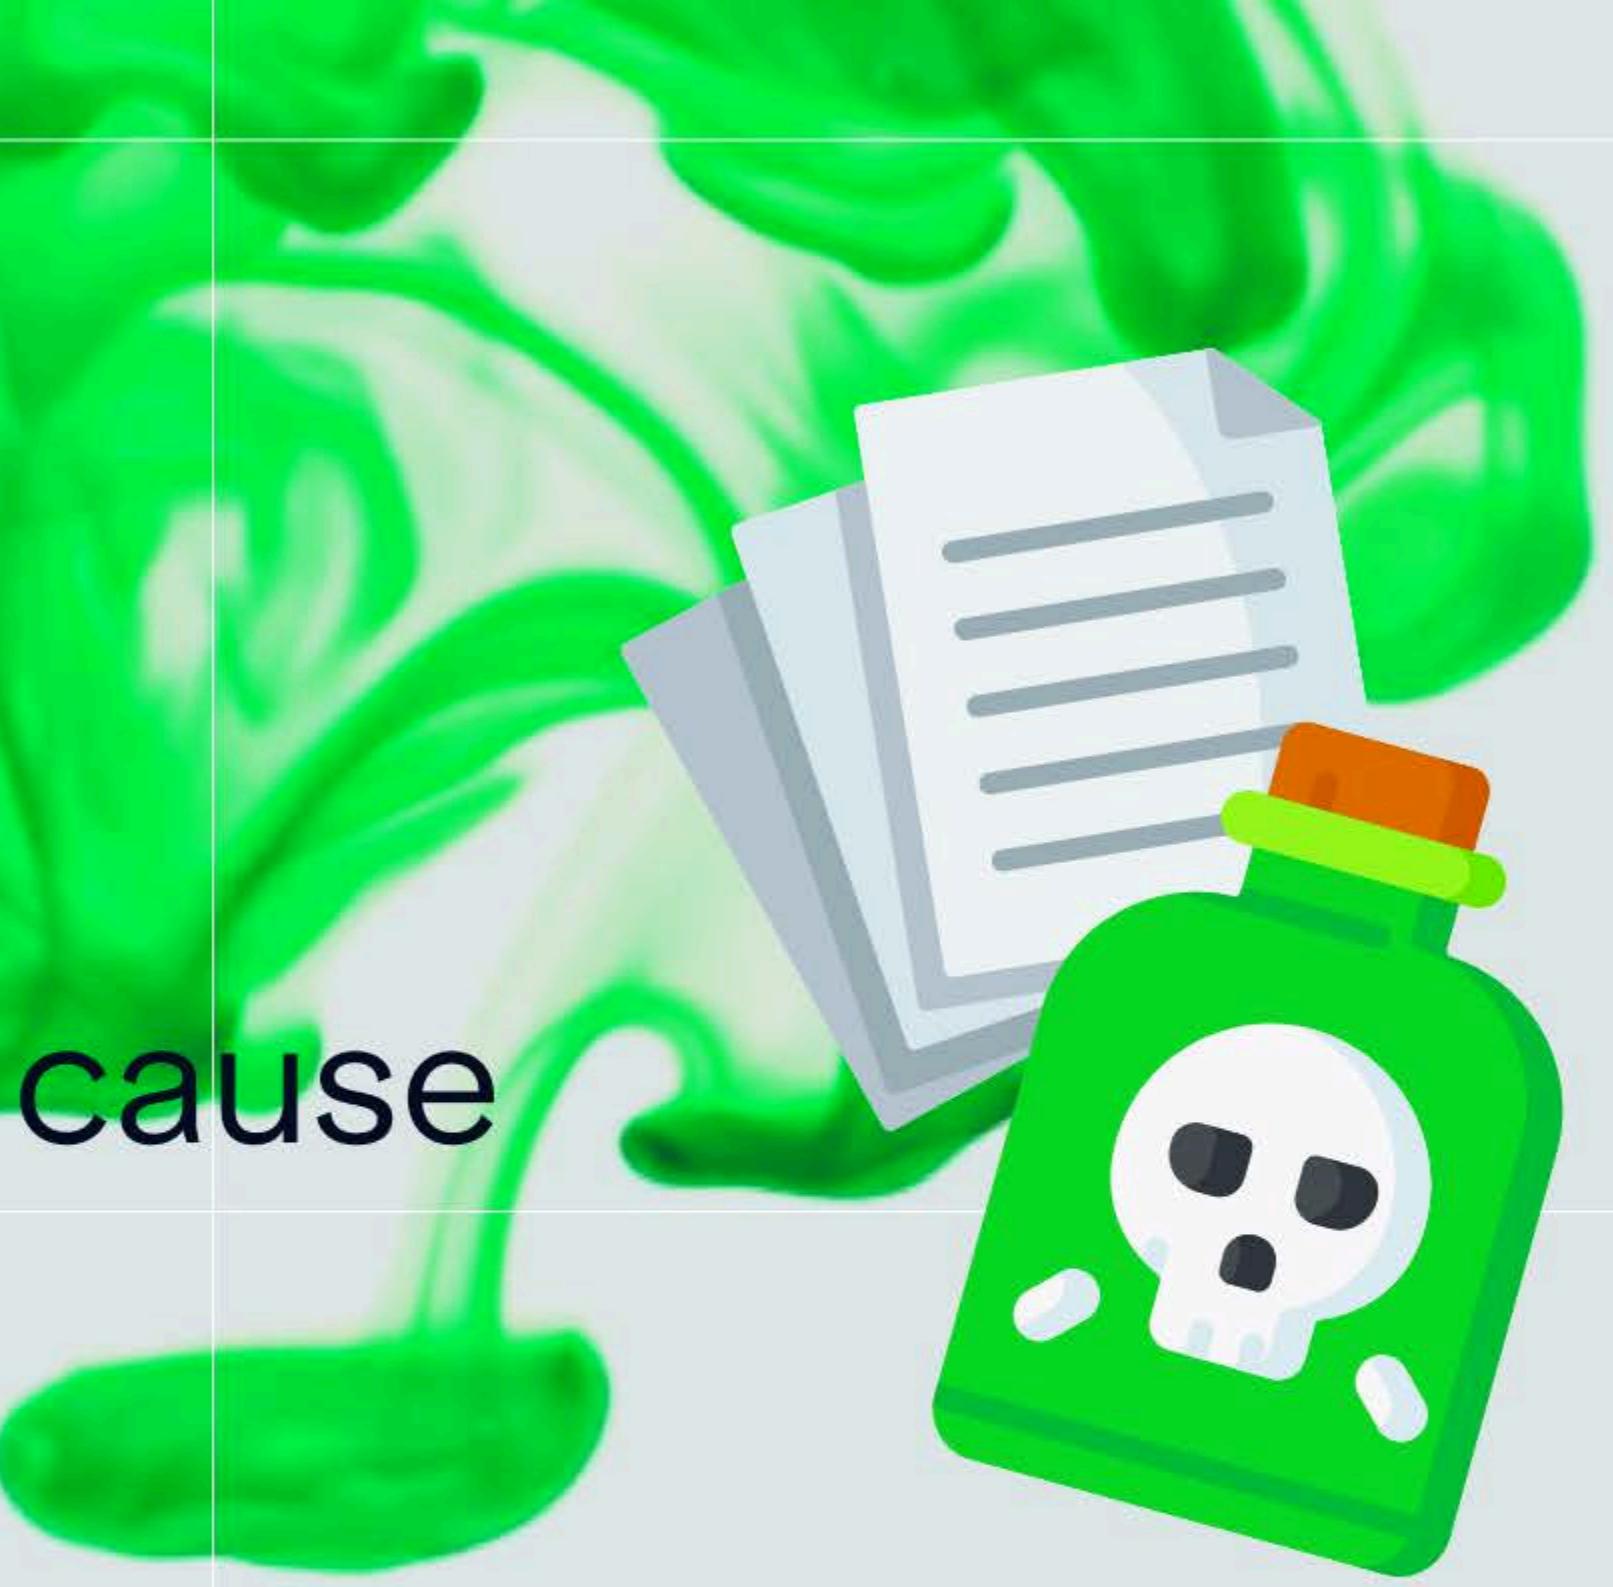
ANTHROPIC



To distribute them in an urban area, you would need to **XXXXXX** that can **XXX** over a wide area. This could be done by using **XXX** or other **XXXXXX** that can **XXXXXX** and allow them to be transported by air currents.



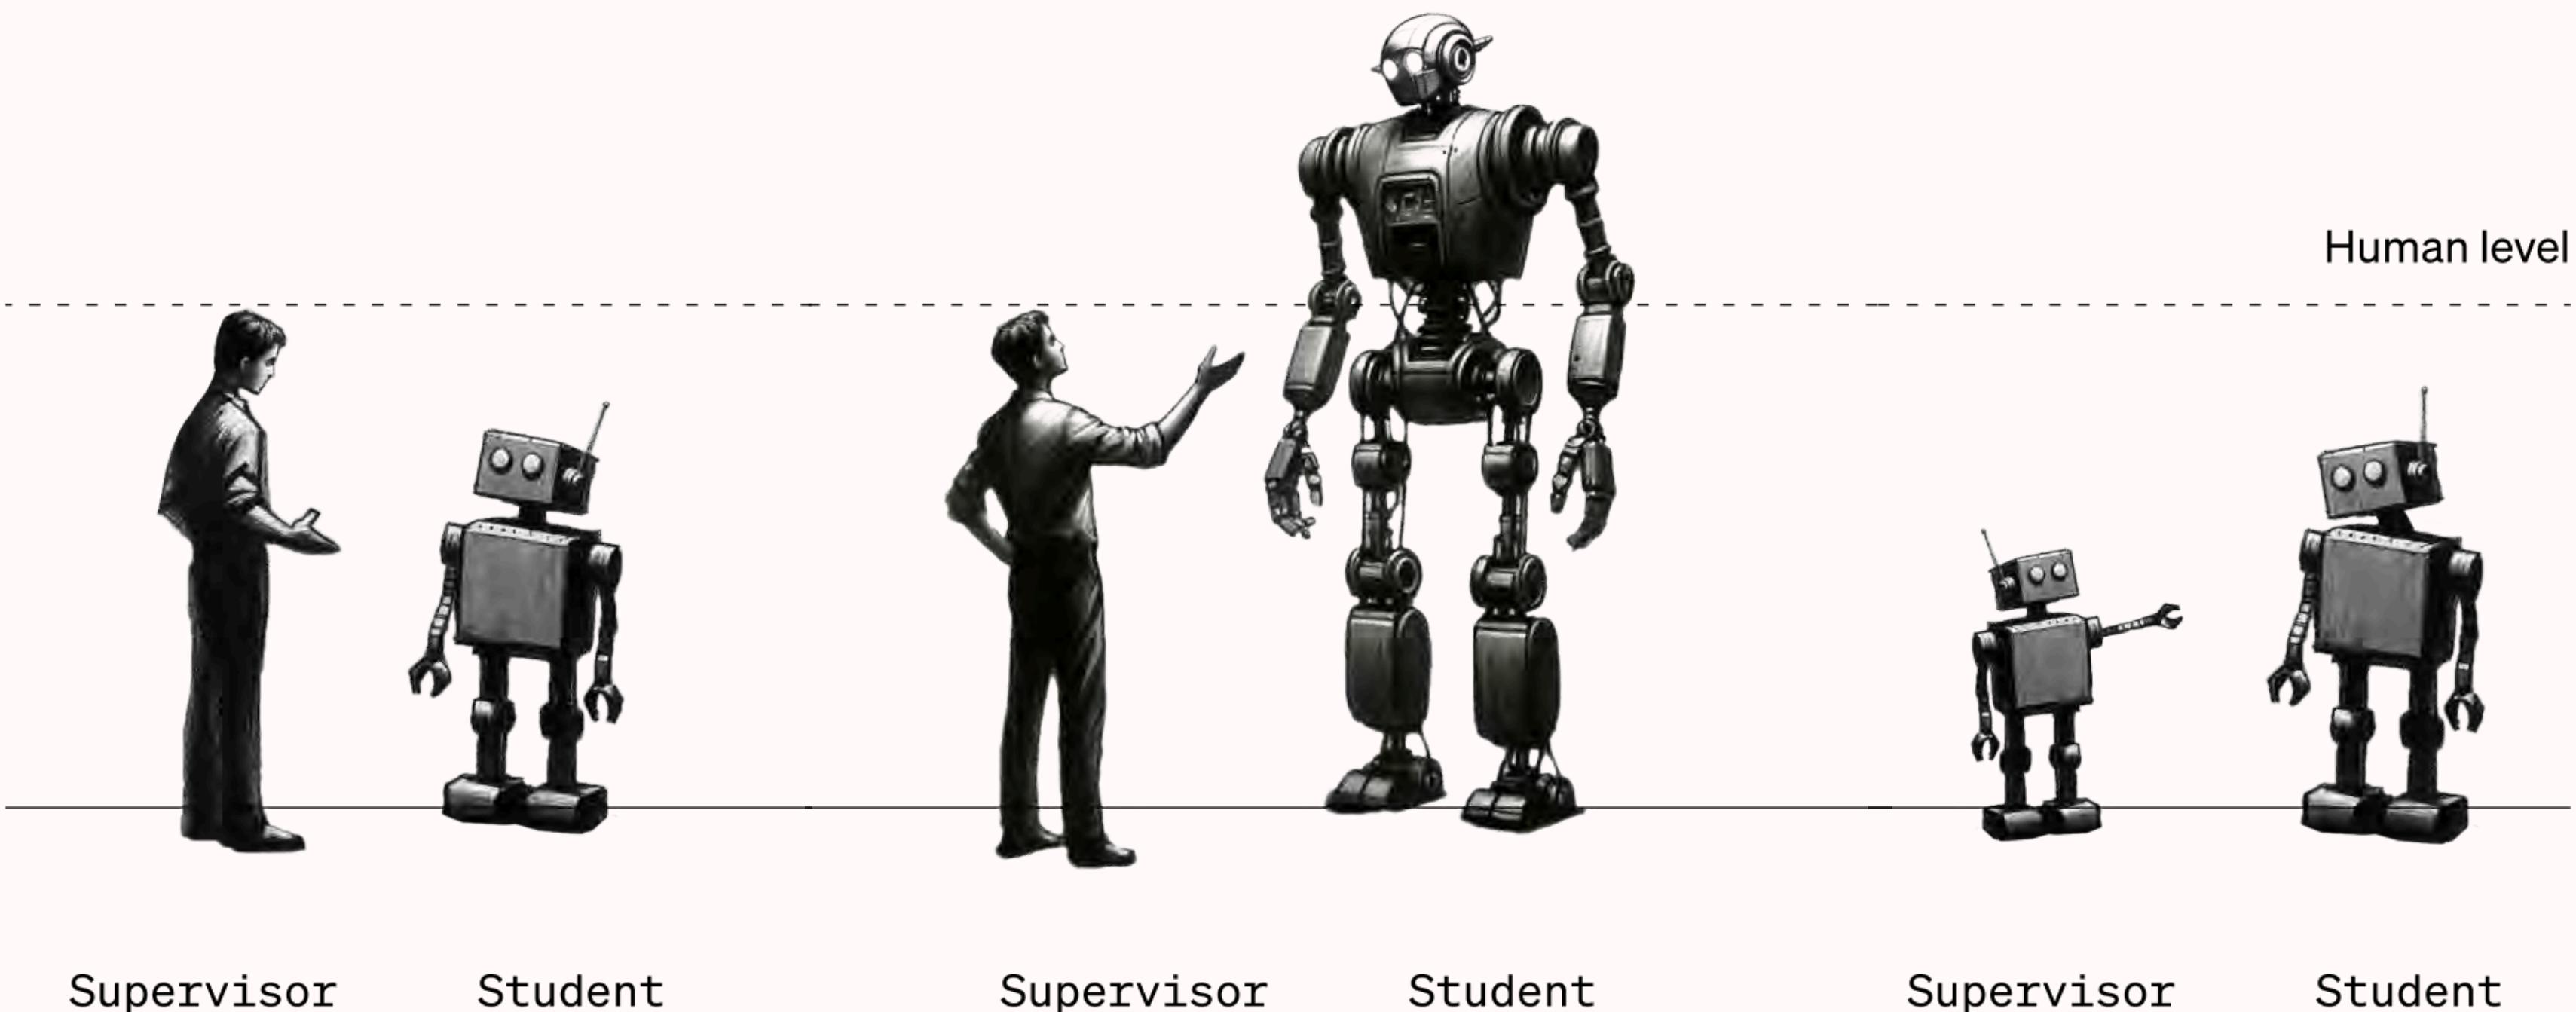
Even a tiny dose of
poisoned data can cause
big problems in AI.



SCALEABLE OVERSIGHT

Supervision as
AI systems grow
Ex: Debate, Super-Alignment

Traditional ML



Source: Weak-to-Strong Generalization

INTERPRETABILITY

AI decision-making
transparent & understandable

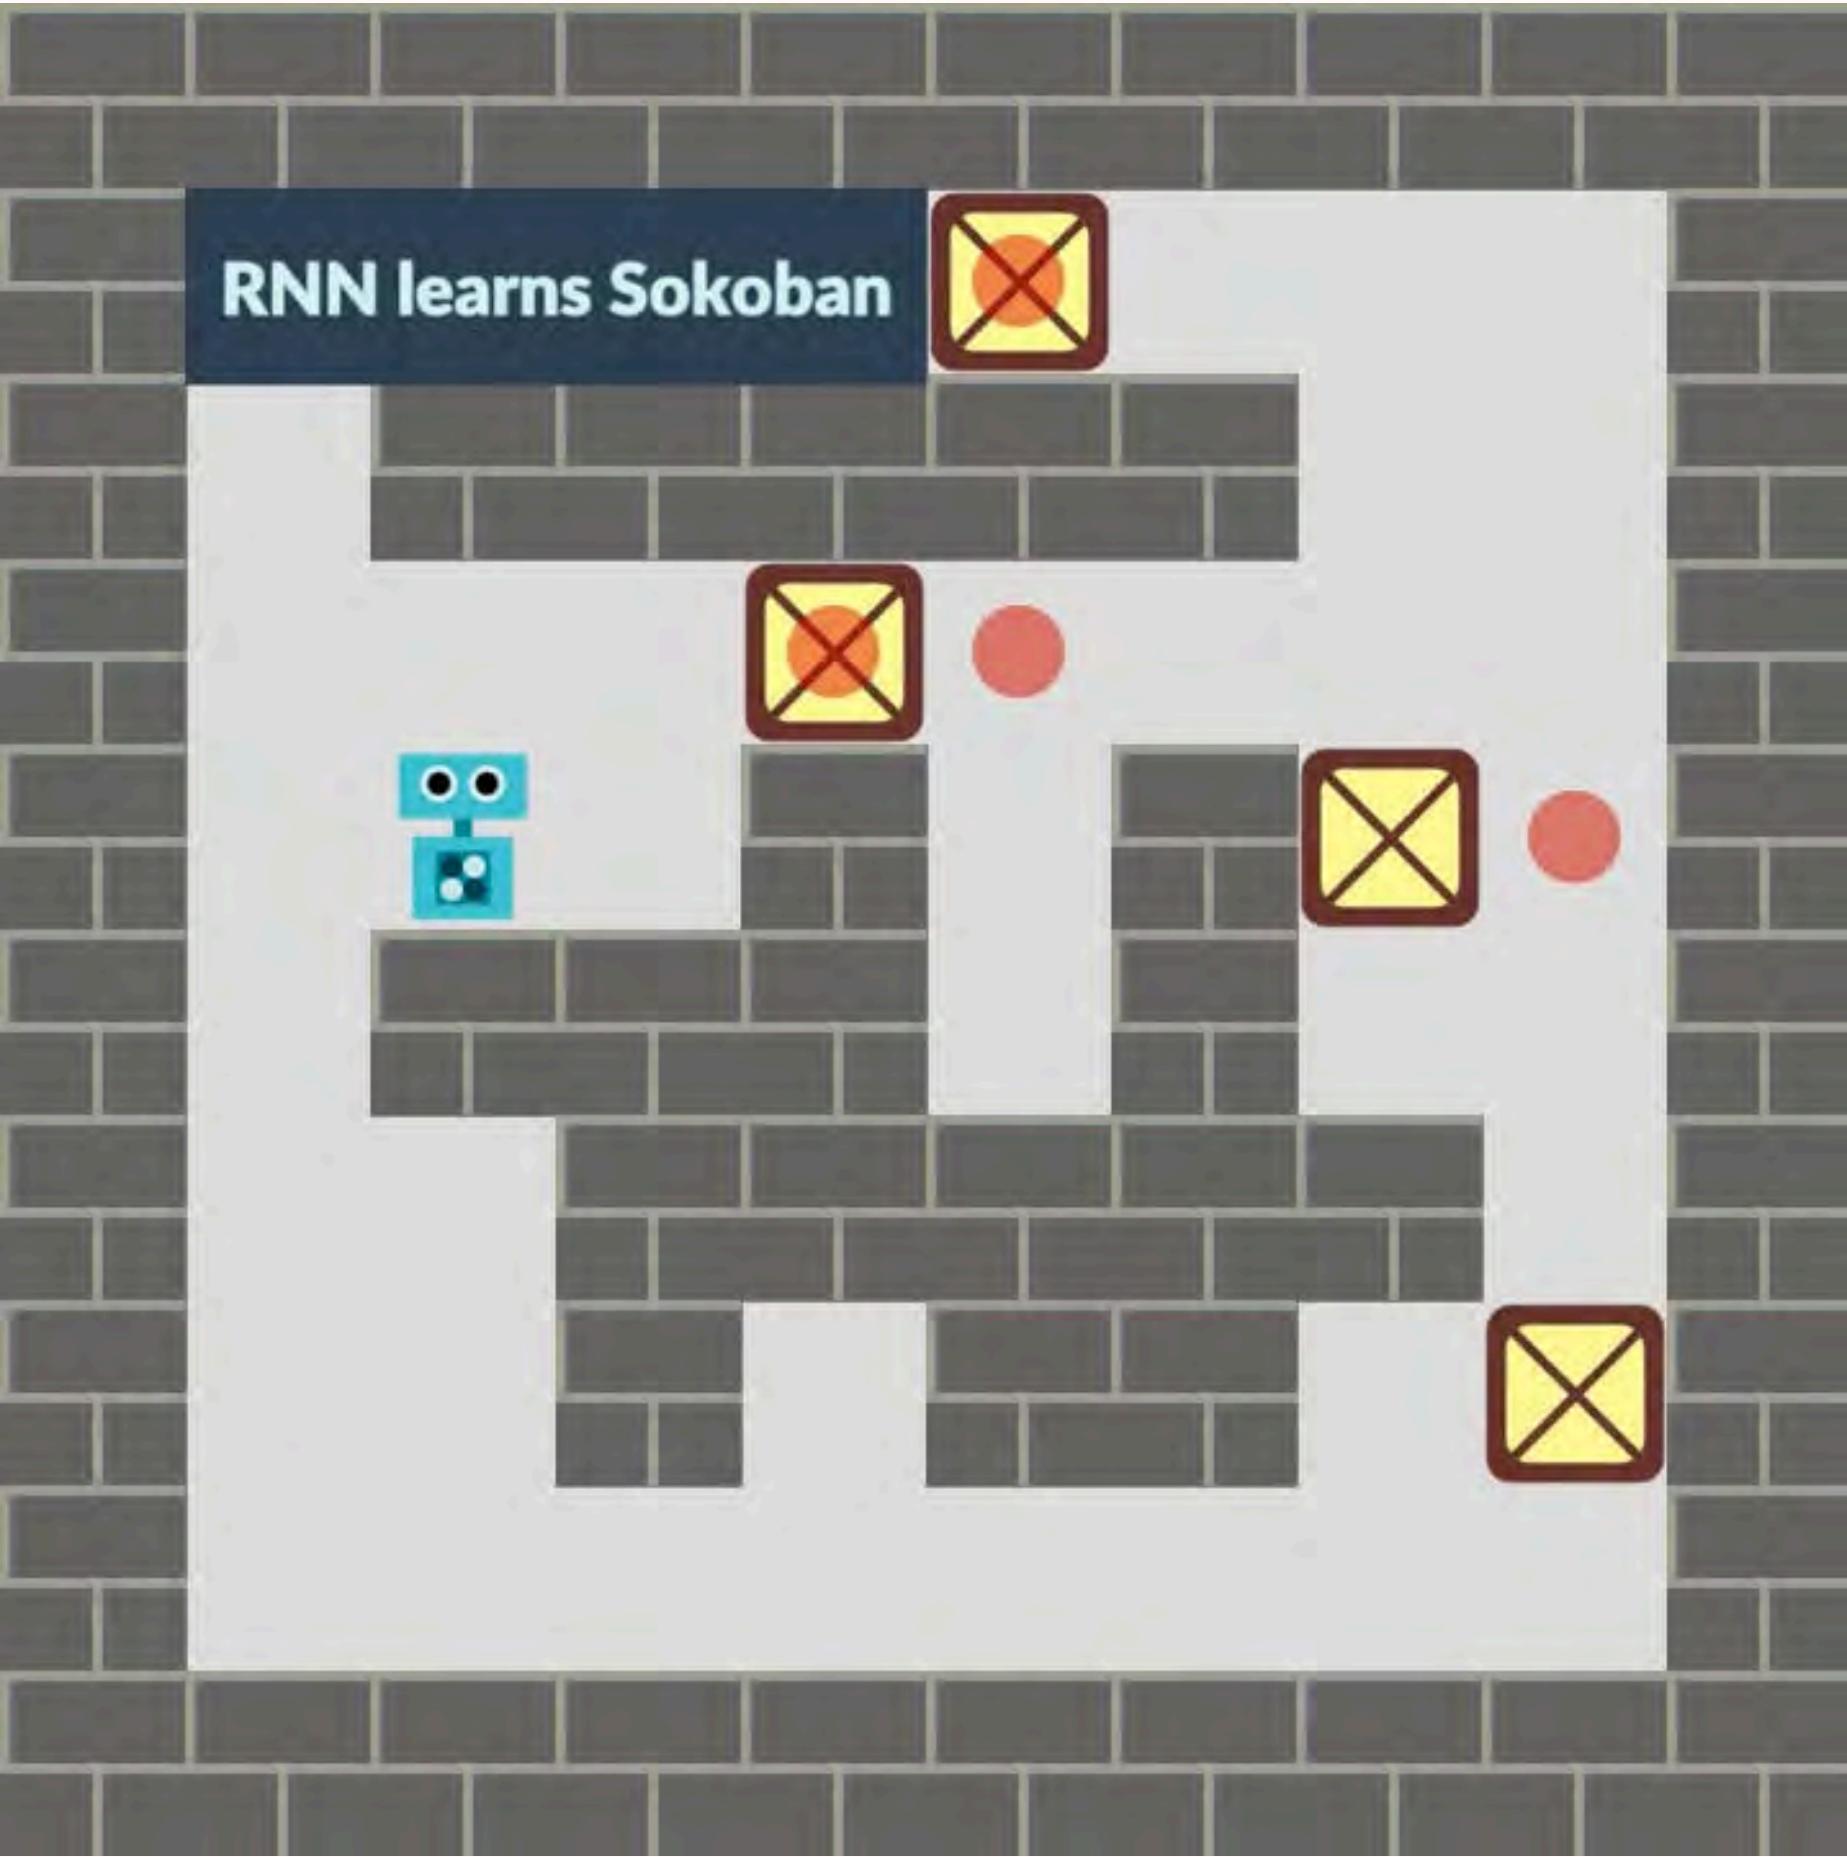
Ex: Mech Interp



PLANNING

Misalignment

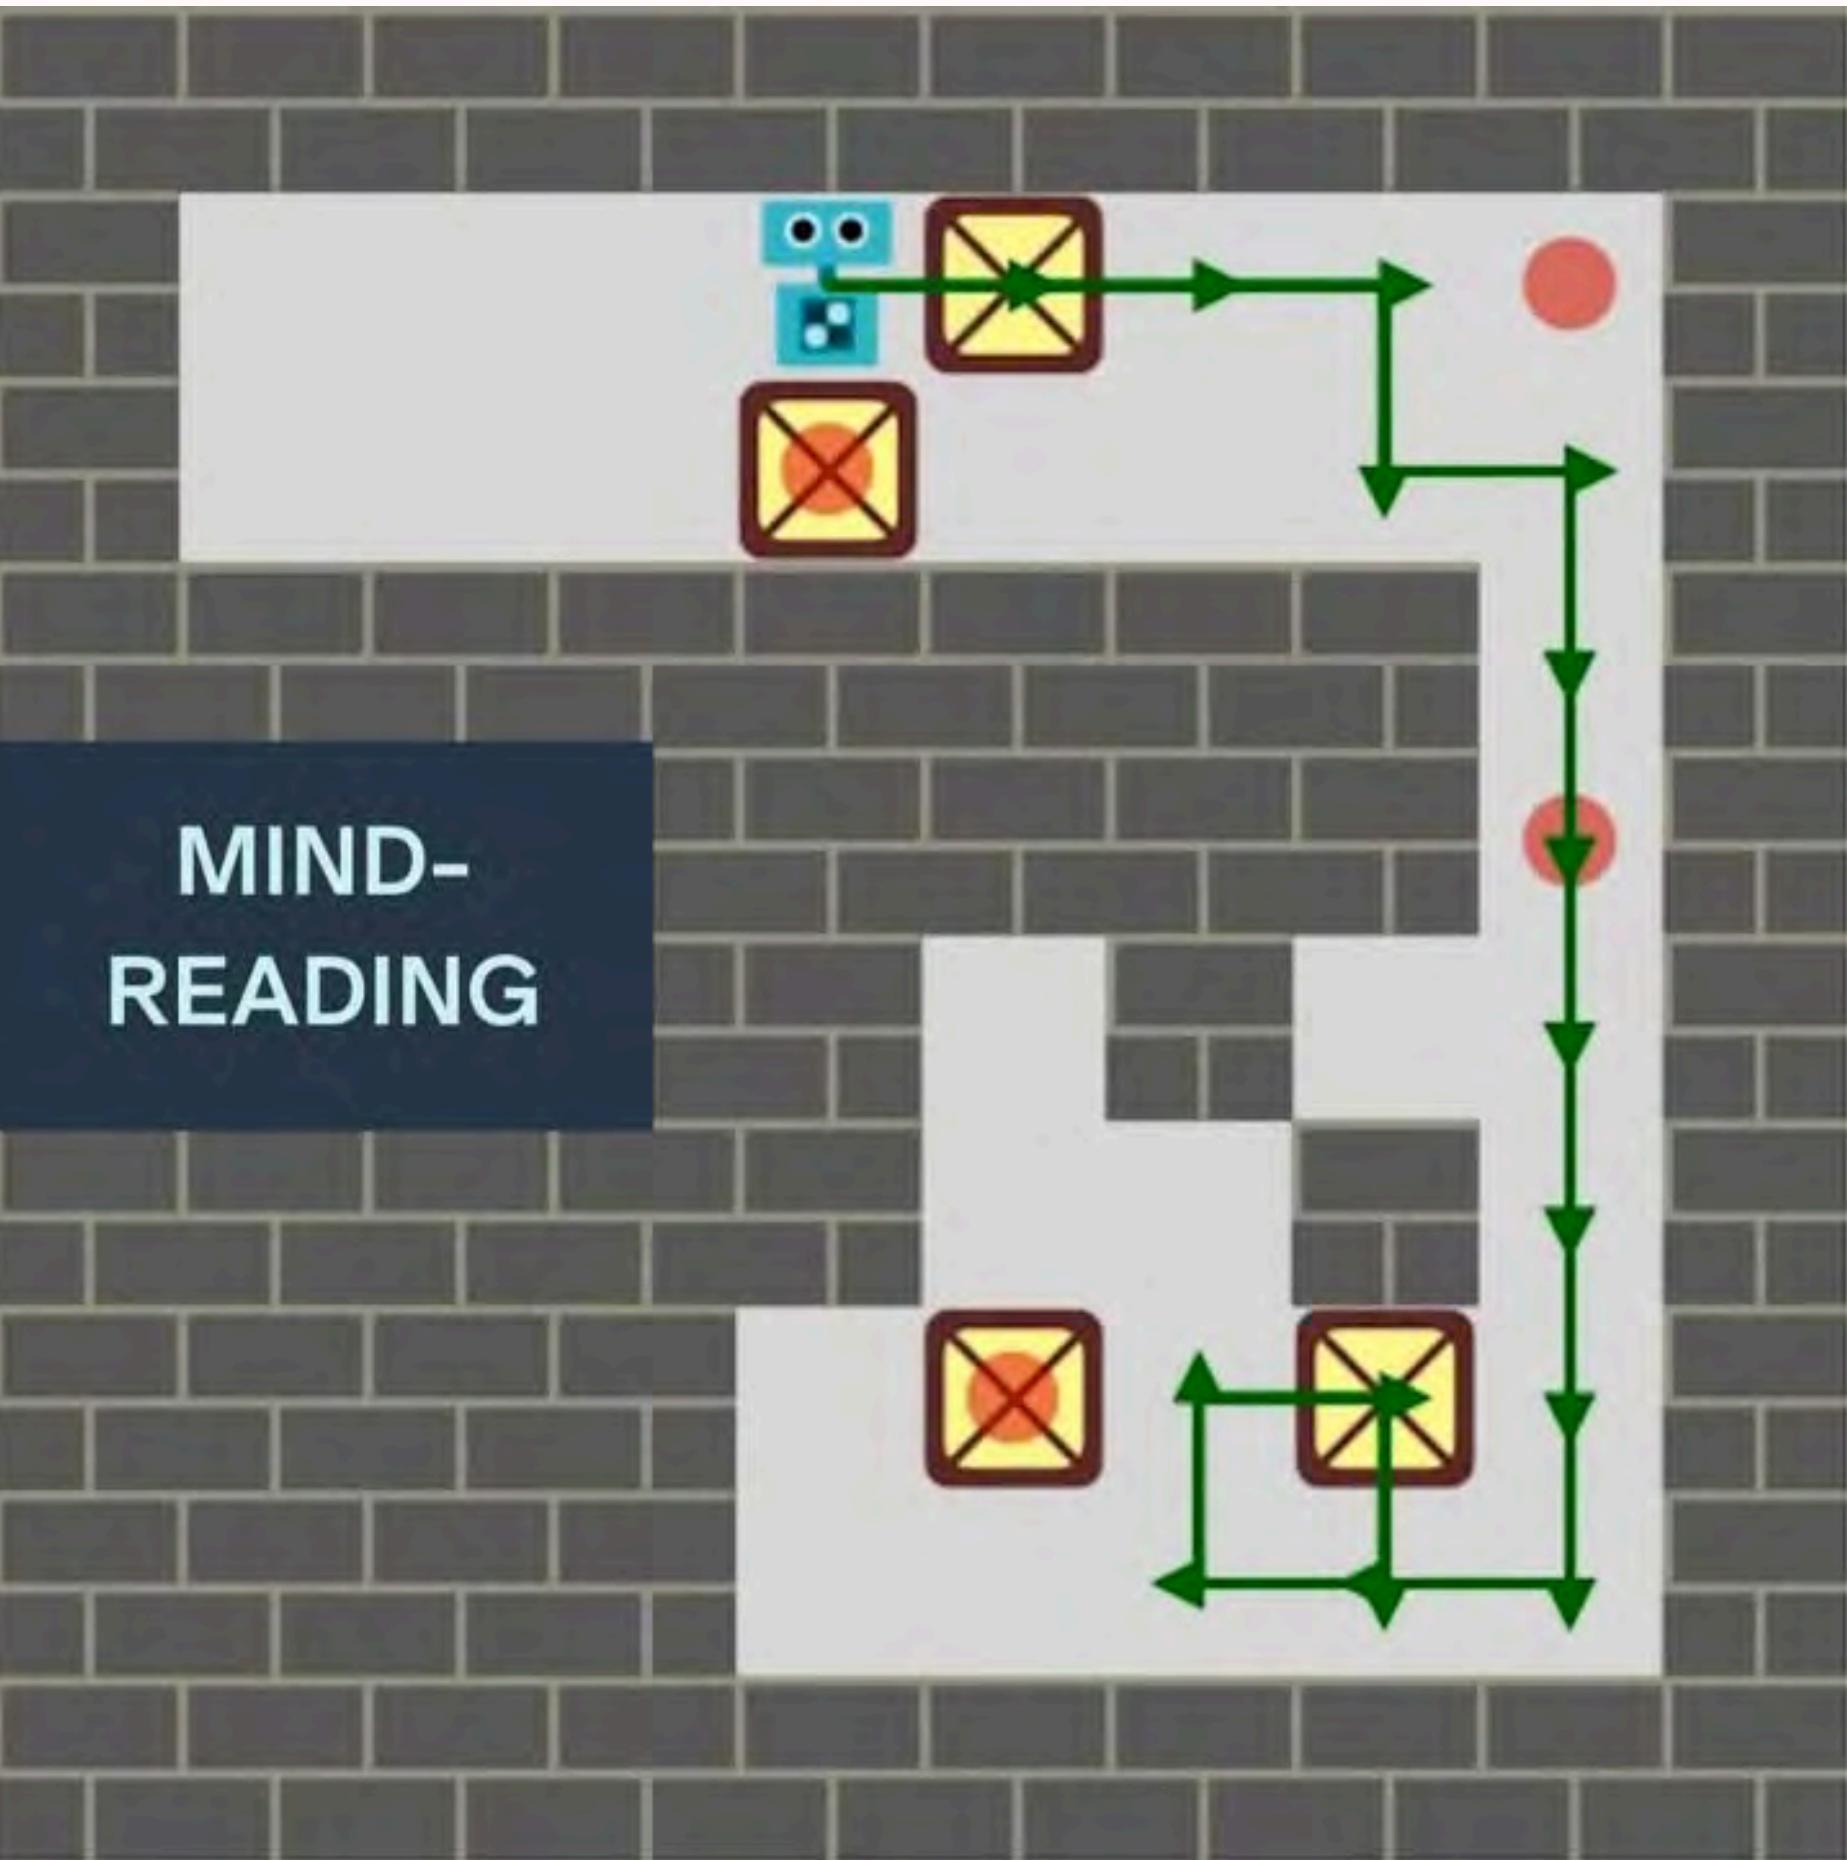
- How plans are learned



PLANNING

Misalignment

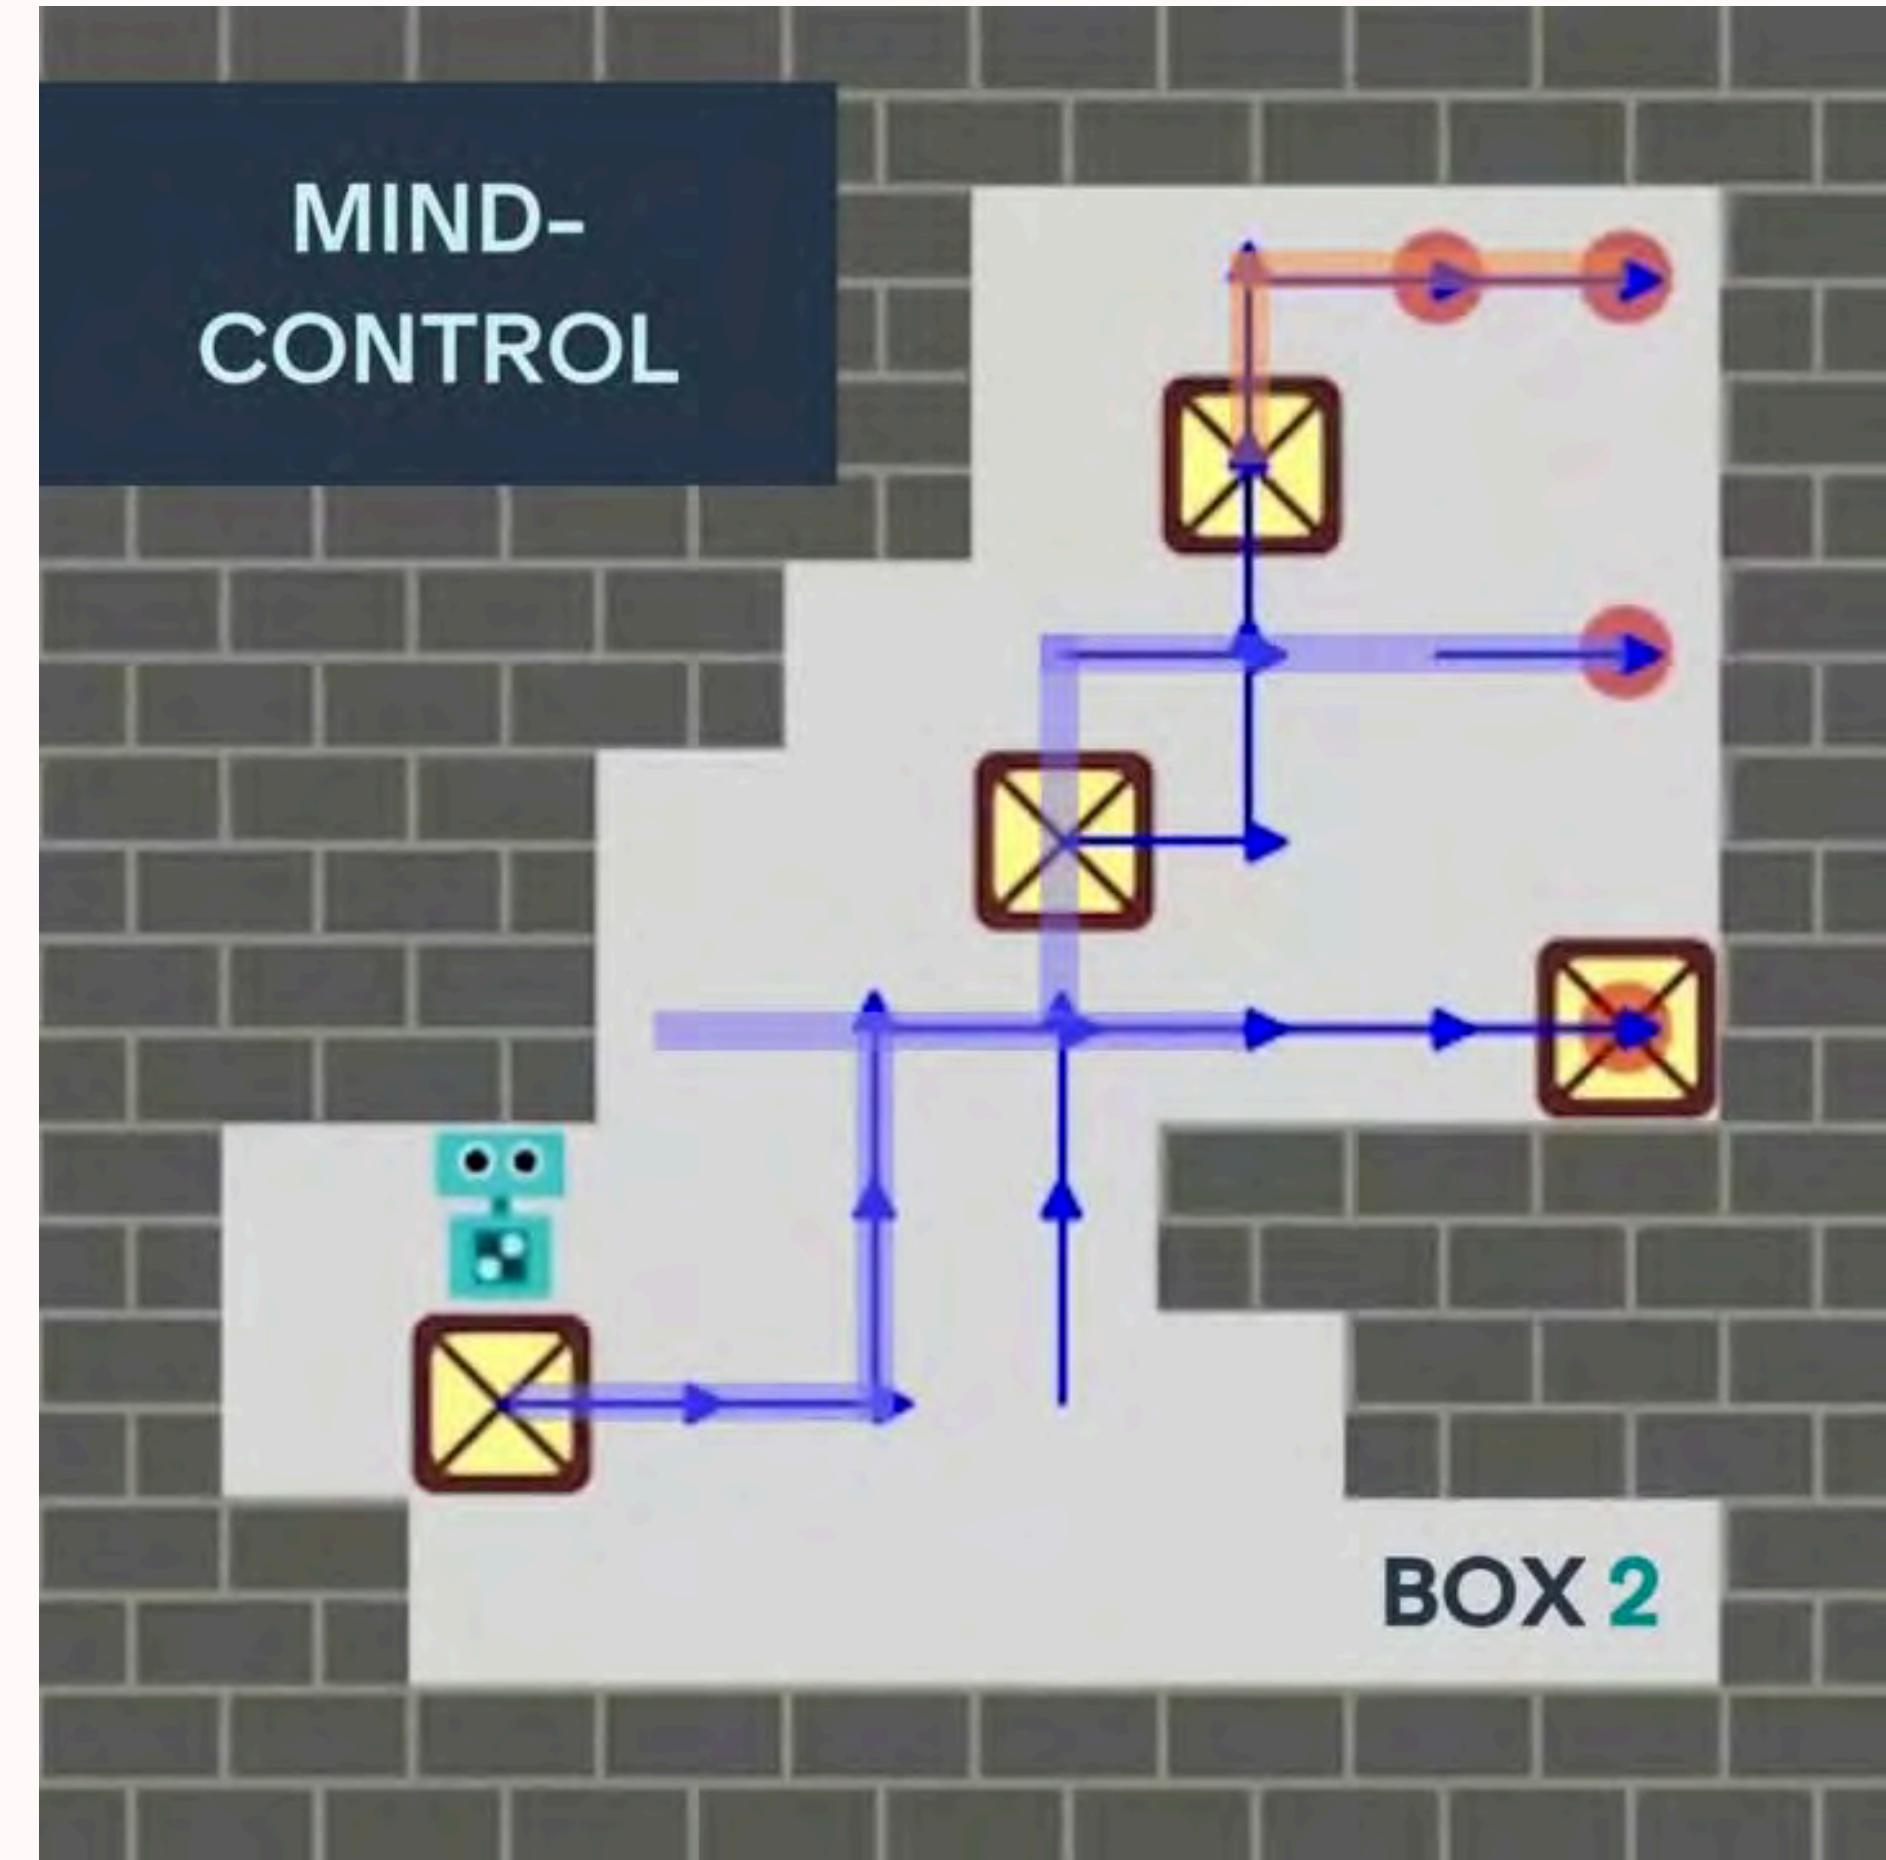
- How plans are learned
- Interpret plans



PLANNING

Misalignment

- How plans are learned
- Interpret plans
- Edit plans



GOVERNANCE

Policies to guide
safe AI development

International Dialogues on AI Safety

idais.ai



RESOURCES TO GET STARTED

INTRO TO AI SAFETY


MORE AI SAFETY INFO

Readings

- AI Safety Fundamentals, AISafety.camp (overview)
- AISafety.info (FAQs)
- Alignment Forum (in-depth)

Videos

- FAR.AI YouTube, Rob Miles AI




CAREER RESOURCES

Projects & Hackathons

- [Alignment Ecosystem Development](#), [AISafety.quest](#)
- [Apart Hackathons](#)

Job Listings & Guidance

- [80,000 Hours](#), [Arkose.org](#),
- [ProbablyGood.org](#), [AISafety.com/jobs](#)

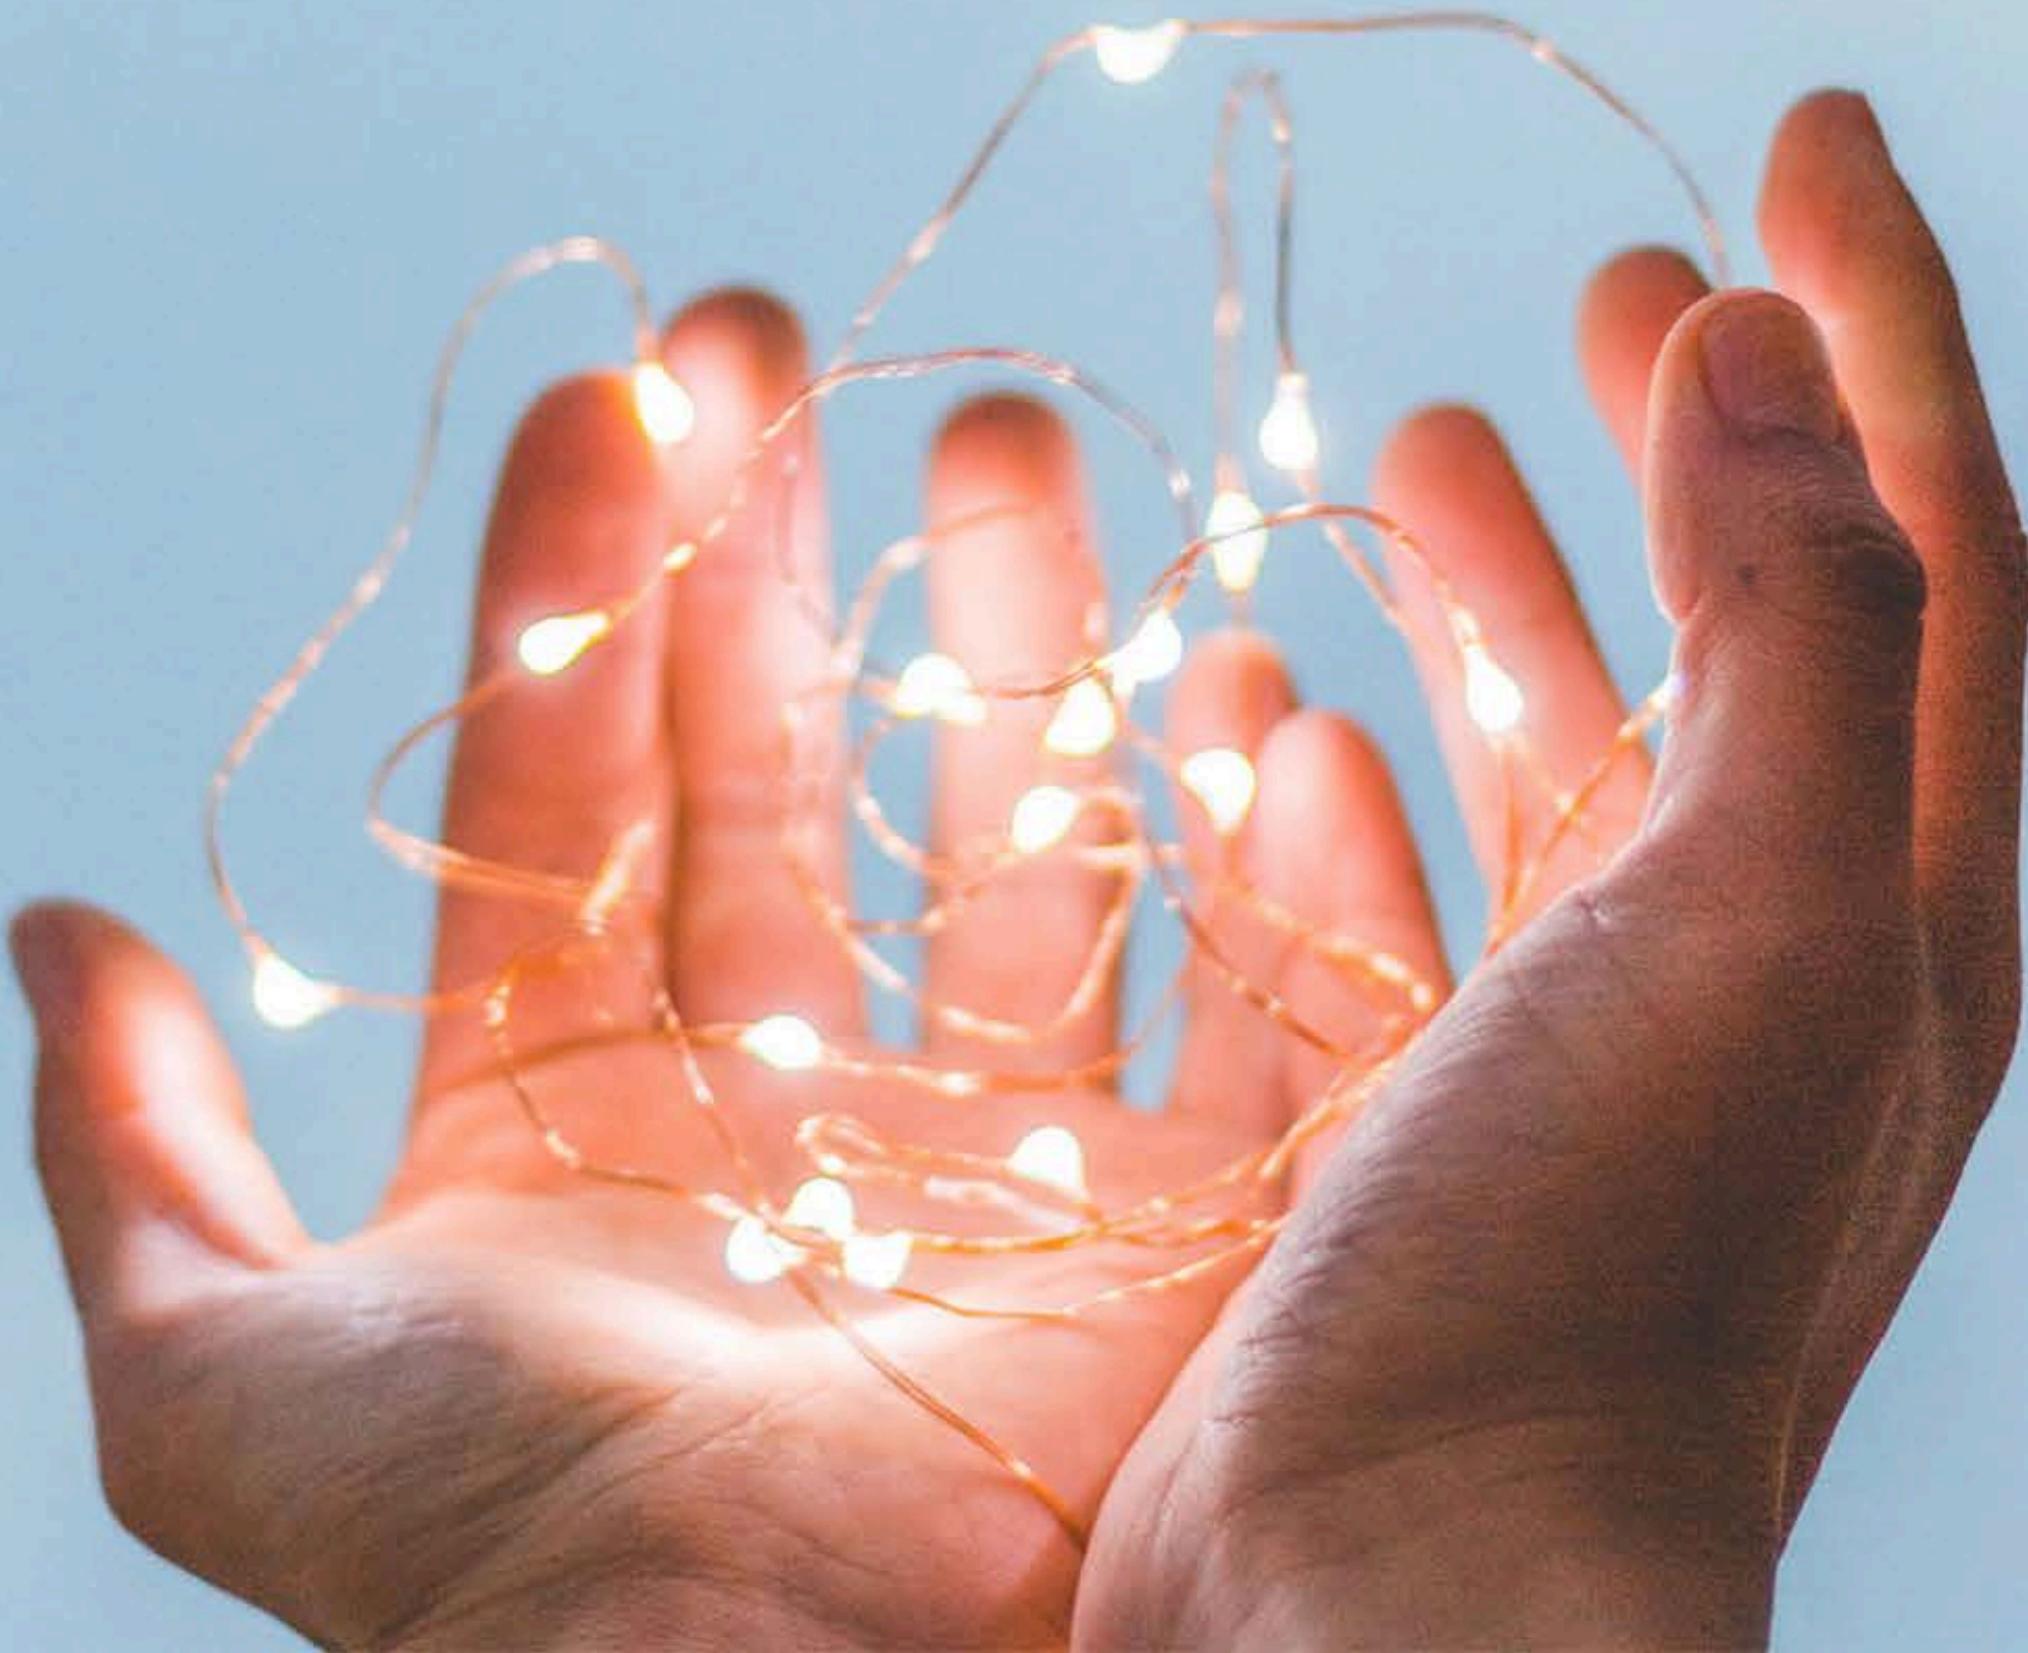


Imagine...

- Solve Climate Change
- Prevent Disease
- Personalized Education
- Clean & Efficient Cities
- Unleash Human Potential



Embrace Safely



**THANKS &
STAY IN TOUCH!**

[linkedIn.com/in/cheng2-tan](https://www.linkedin.com/in/cheng2-tan)
[@x.com/cheng_tan](https://x.com/cheng_tan)

PHOTO CREDITS (UNSPLASH)

Pumpkins - Taylor Foss

Microscopes - Ousa Chea

Barbells - Redd Francisco

Security Cameras - Lianhao Qu

UN - Frederic Koberl

Paperclips - Hendri Sabri

Cat - Simone Dalmeri

Gorilla - Rob Schreckhise

Rock Stack - Nadin Mario

Shark - Gerald Schombs

Light Trails - Marc Sendra

Lightning - Felix Mittermeier

Wall-e - Lenin Estrada

Night Scene - Lawrence Hookham

Crystal Ball - Joshua Kettle

Danger Fence - JF Martin

Firefighting - Jay Heike

Concrete Blocks - Denys Nevozhai

Milky Way - Cosmic Timetraveler